



Check 2 — Cross-Subject Replication (LoRA mid4★, best model)

Does the encoding model capture a generalizable language-to-brain mapping?

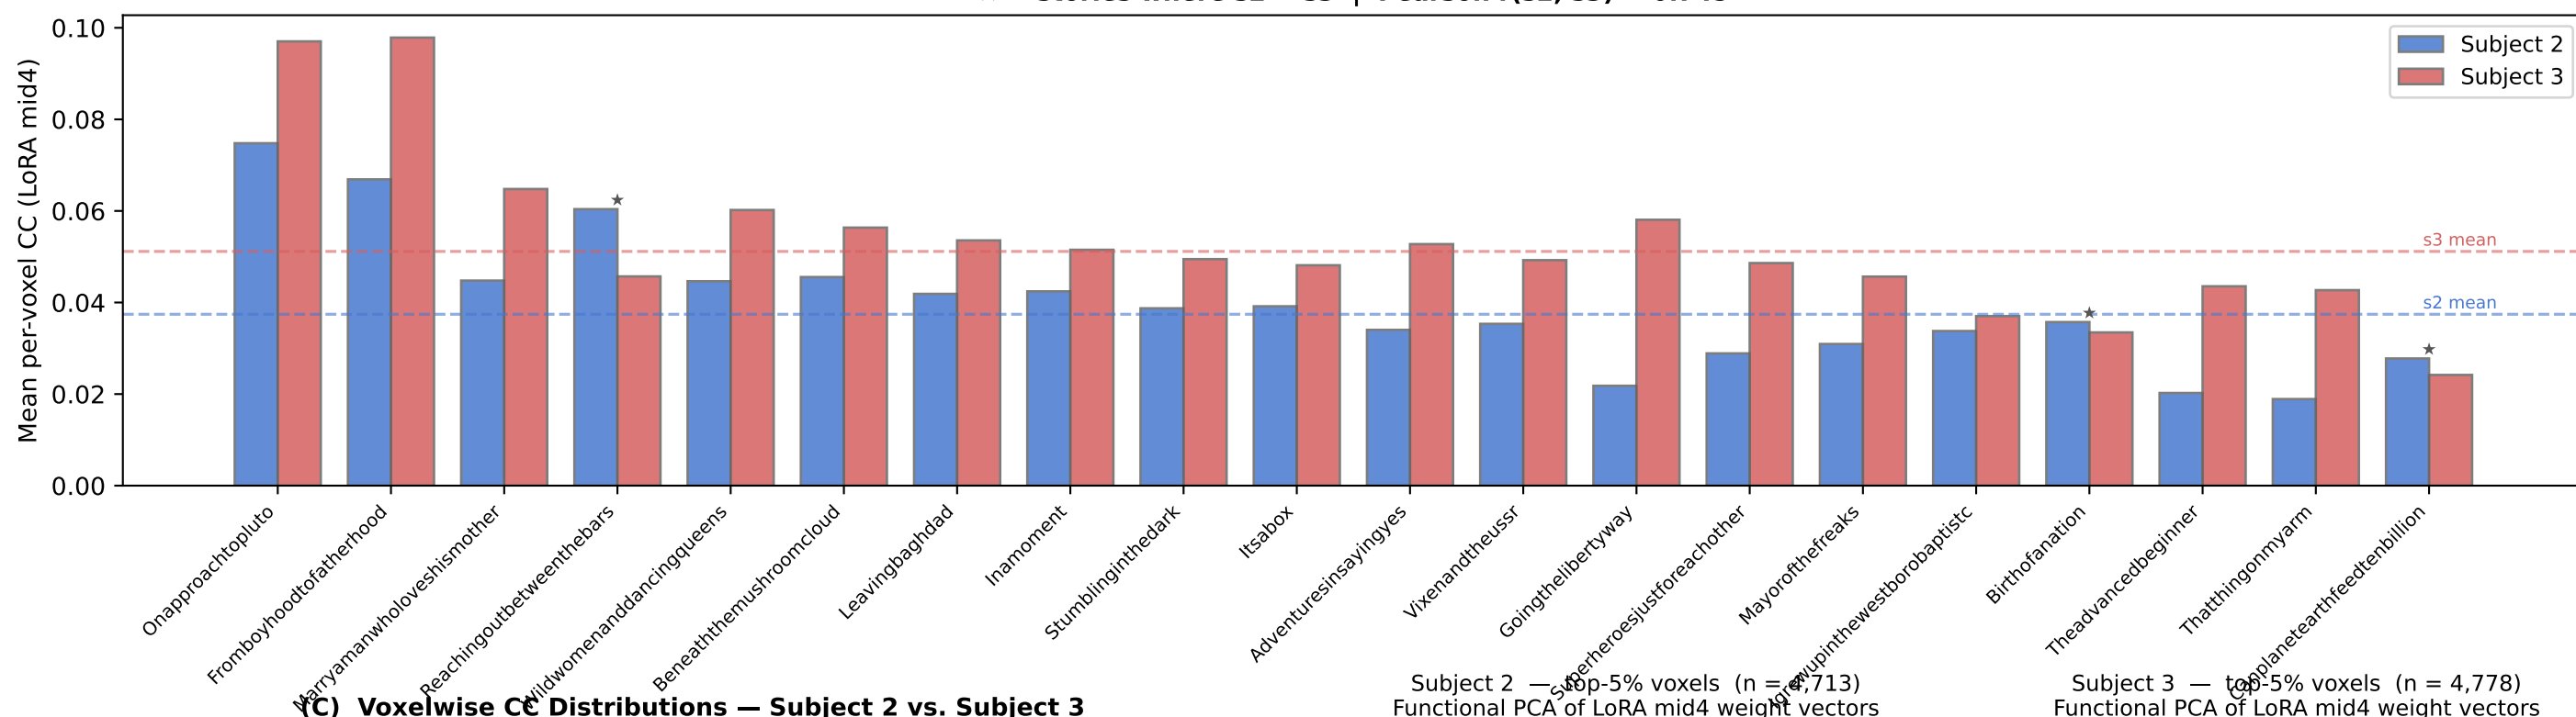
(A) Story-level Encoding Agreement
Pearson $r = 0.745$, $n = 20$ test stories



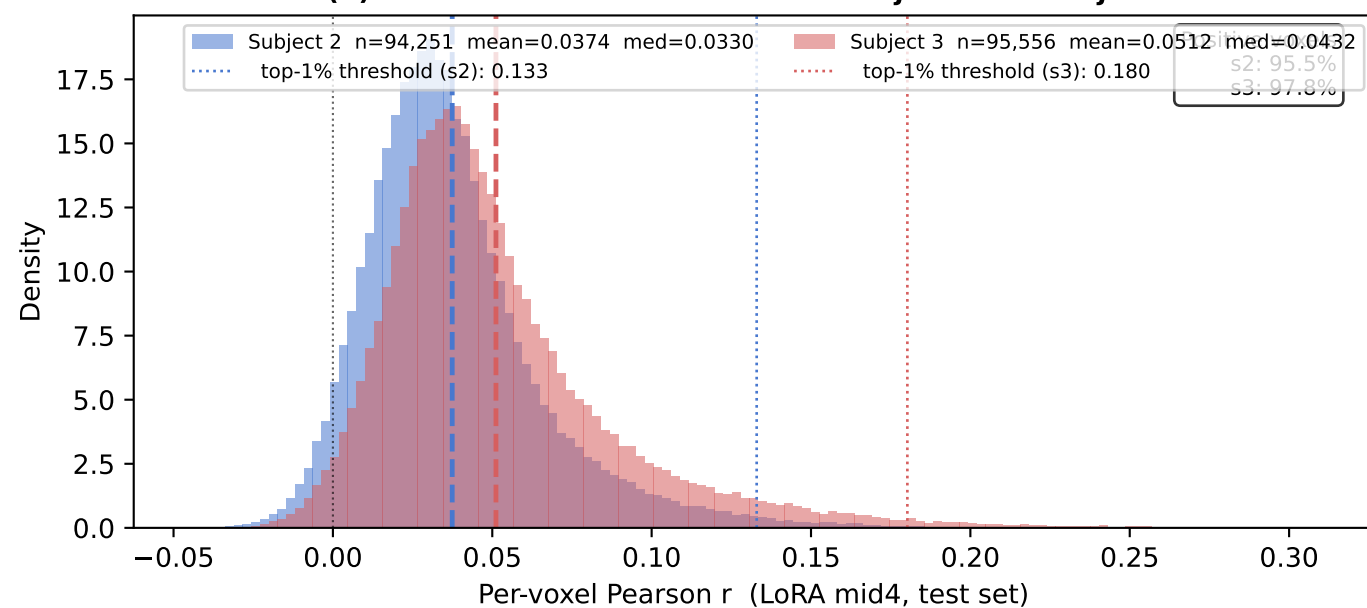
(E) Cross-Subject CC by Embedding Method
 $s3/s2$ ratio ($\times$) shrinks as models improve: subjects converge



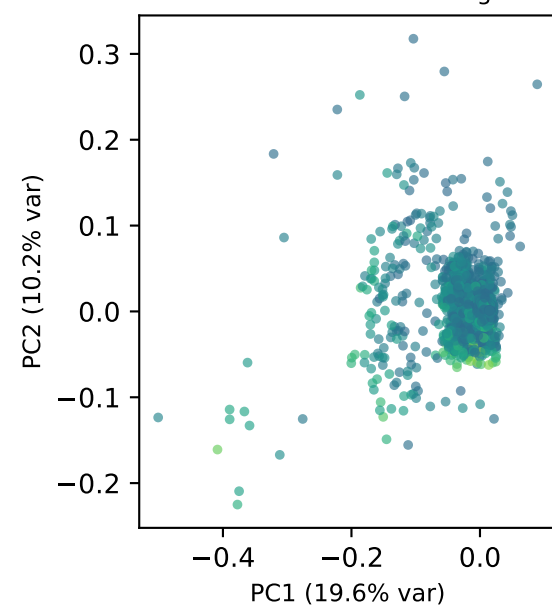
(B) Per-Story Encoding Accuracy — Both Subjects (sorted by mean CC)
★ = stories where $s2 > s3$ | Pearson $r(s2, s3) = 0.745$



(C) Voxelwise CC Distributions — Subject 2 vs. Subject 3



Subject 2 — top-5% voxels ($n = 4,713$)
Functional PCA of LoRA mid4 weight vectors



Subject 3 — top-5% voxels ($n = 4,778$)
Functional PCA of LoRA mid4 weight vectors

